


Kubernetes 101 – an introduction



Gerlof Langeveld

`www.atcomputing.nl`




```
git clone https://github.com/atcomputing/k8soverview
```

Container intro – conventional ecosystem

Conventional production environment


- all applications use same 'ecosystem'
 - filesystem
 - network
 - process numbers (PIDs)
 - user identities
 - hardware resources
- disadvantages
 - laborious (de)installation of application
 - compromised application might
 - access all data
 - manipulate entire network stack
 - compromise other applications
 - application might overload resources



Container intro – isolated ecosystem

Container: isolated ecosystem for application


- own mini filesystem containing
 - program executables
 - libraries
 - configuration files
 - data files
 - ...
- own network
- own PID numbers
- own user identities
- cpu/memory/disk guarantees and restrictions



Container intro – container image

Image

- needed to start container on destination host
- contains
 - mini filesystem
 - metadata, e.g. application to be started
- images stored in *registry*
- *developer* ('dev')
 - builds image for application
 - pushes image to registry
- *operations* ('ops')
 - pulls image from registry
 - uses image to activate container



Container intro – build custom image

Build custom image Docker example

- specify base image and own modifications in file **Dockerfile**

```
$ cat Dockerfile
FROM ubuntu:18.04
RUN apt-get update && apt-get install -y apache2
COPY index.html /var/www/html/index.html
CMD ["/usr/sbin/apache2ctl", "-D", "FOREGROUND"]
```

```
$ cat index.html
<h1> Message from container! </h1>
```

- build custom image

```
$ docker build -t atcomp/apachetest .
Successfully built 5d3b567581df
```

new image

directory containing **Dockerfile**
and other files needed in image

- list images

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
atcomp/apachetest	latest	5d3b567581df	About an hour ago	204MB
ubuntu	18.04	5d2df19066ac	3 weeks ago	63.1MB

Container intro – run from custom image

Use custom image to run container Docker example

- run custom container

```
$ docker run -p 8080:80 -d atcomp/apachetest
```

publish port *detached: run in background*

- contact webserver via URL `http://localhost:8080`
by web browser or command `curl`

```
$ curl http://localhost:8080  
<h1> Message from container! </h1>
```

- push image to registry

```
$ docker push atcomp/apachetest
```

Kubernetes 101 – an introduction



Kubernetes
basics

What is Kubernetes?

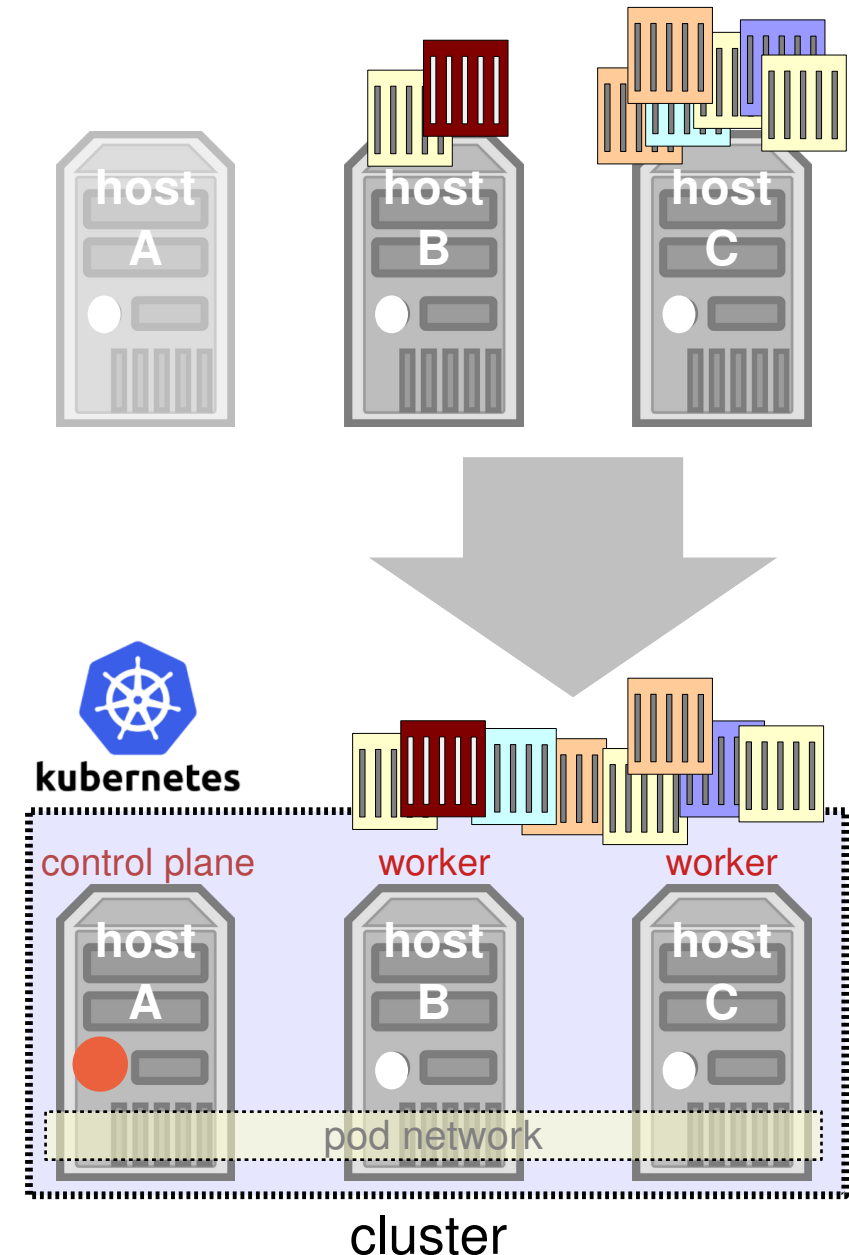
Kubernetes ('helmsman') aka. K8s

- combines various hosts into *cluster*
 - scalability
 - reliability (failover)
- orchestrates containers
 - activate
 - monitor
 - terminate
- uses container runtime implementation, like *containerd*, *cri-o*,

Kubernetes orchestration

Kubernetes orchestration

- introduced in 2015,
inspired by Google's Borg
- maintained by
Cloud Native Computing Foundation (CNCF)
- concept
 - cluster needs at least one *control-plane node*,
formerly known as *master node*
 - other hosts in cluster are *worker nodes*
that only run container instances



User interface

User interfaces

- command line interface: **kubectl subcommand object [options]**

- subcommands

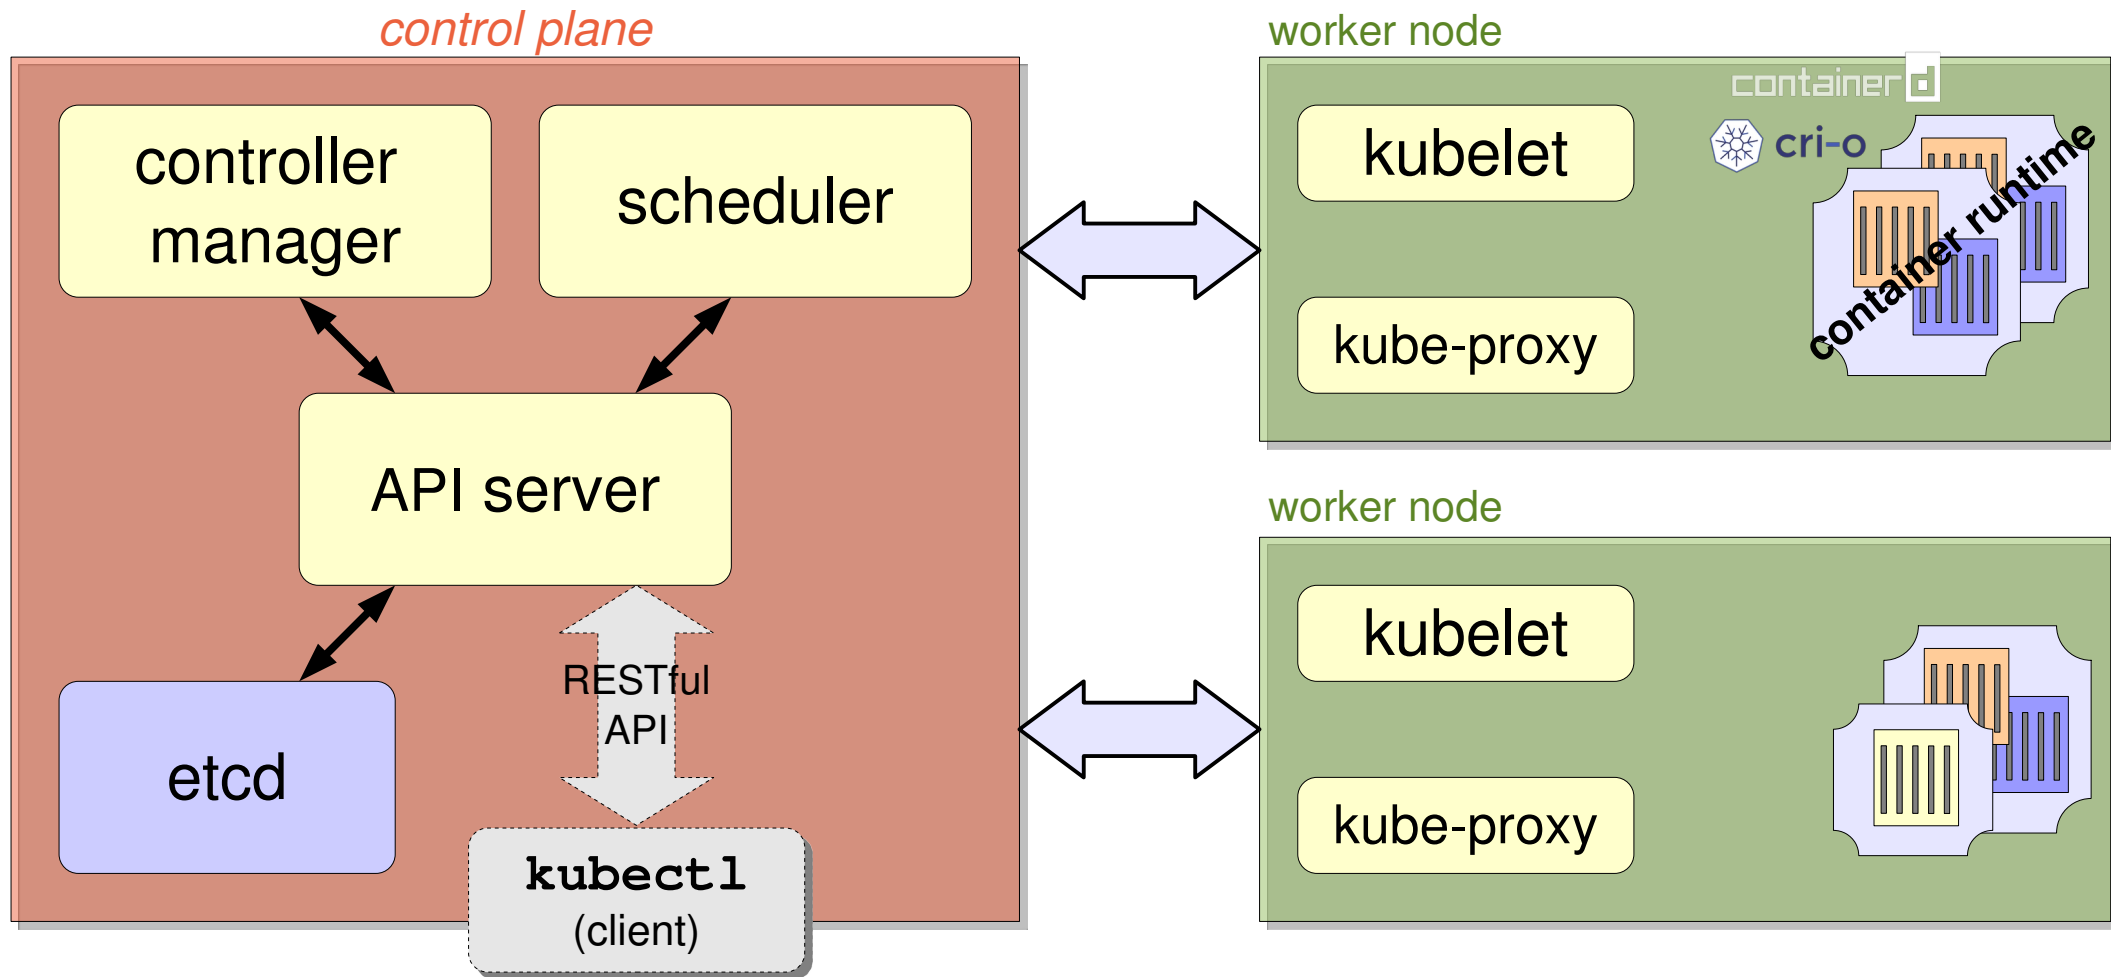
- create, modify and delete objects: **create, apply, delete**
- query current state of objects: **get, describe**
- other actions, like: **logs, scale,**

```
$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
hosta	Ready	control-plane	75m	v1.26.0
hostb	Ready	<none>	40m	v1.26.0
hostc	Ready	<none>	40m	v1.26.0

- graphical user interface
 - also valid for cloud implementations, like
 - Google Kubernetes Engine (GKE)
 - Amazon Elastic Kubernetes Service (EKS)
 - Azure Kubernetes Service (AKS)

Kubernetes architecture




- **API server:** controls entire cluster
- **scheduler:** schedules containers (via pods) on nodes
- **controller manager:** controls required number of replicas
- **etcd:** distributed key-value store to maintain current cluster state

- **kubelet:**
 - pod startup & monitoring
 - resource management
- **kube-proxy:** container exposure to network & load balancing

Pods – introduction

Pod



- conceptually smallest unit, wrapping one or more containers
 - typically one container running primary application, with other supporting containers
 - containers in pod share same IP address on *internal* pod network
- run pod manually

```
$ kubectl run apatest --restart=Never --image=atcomp/apachetest --port 80
pod/apatest created
```

- or run pod via *manifest* file (YAML syntax)

```
$ kubectl apply -f apa.yml
pod/apatest created
```

- status:

```
$ kubectl get pods
NAME        READY   STATUS    RESTARTS   AGE
apatest     1/1     Running   0           8s

$ kubectl get pods -o wide
...  AGE    IP           NODE    ....
      8s    10.244.1.5   hostb

$ curl 10.244.1.5
<h1> Message from container! </h1>
```

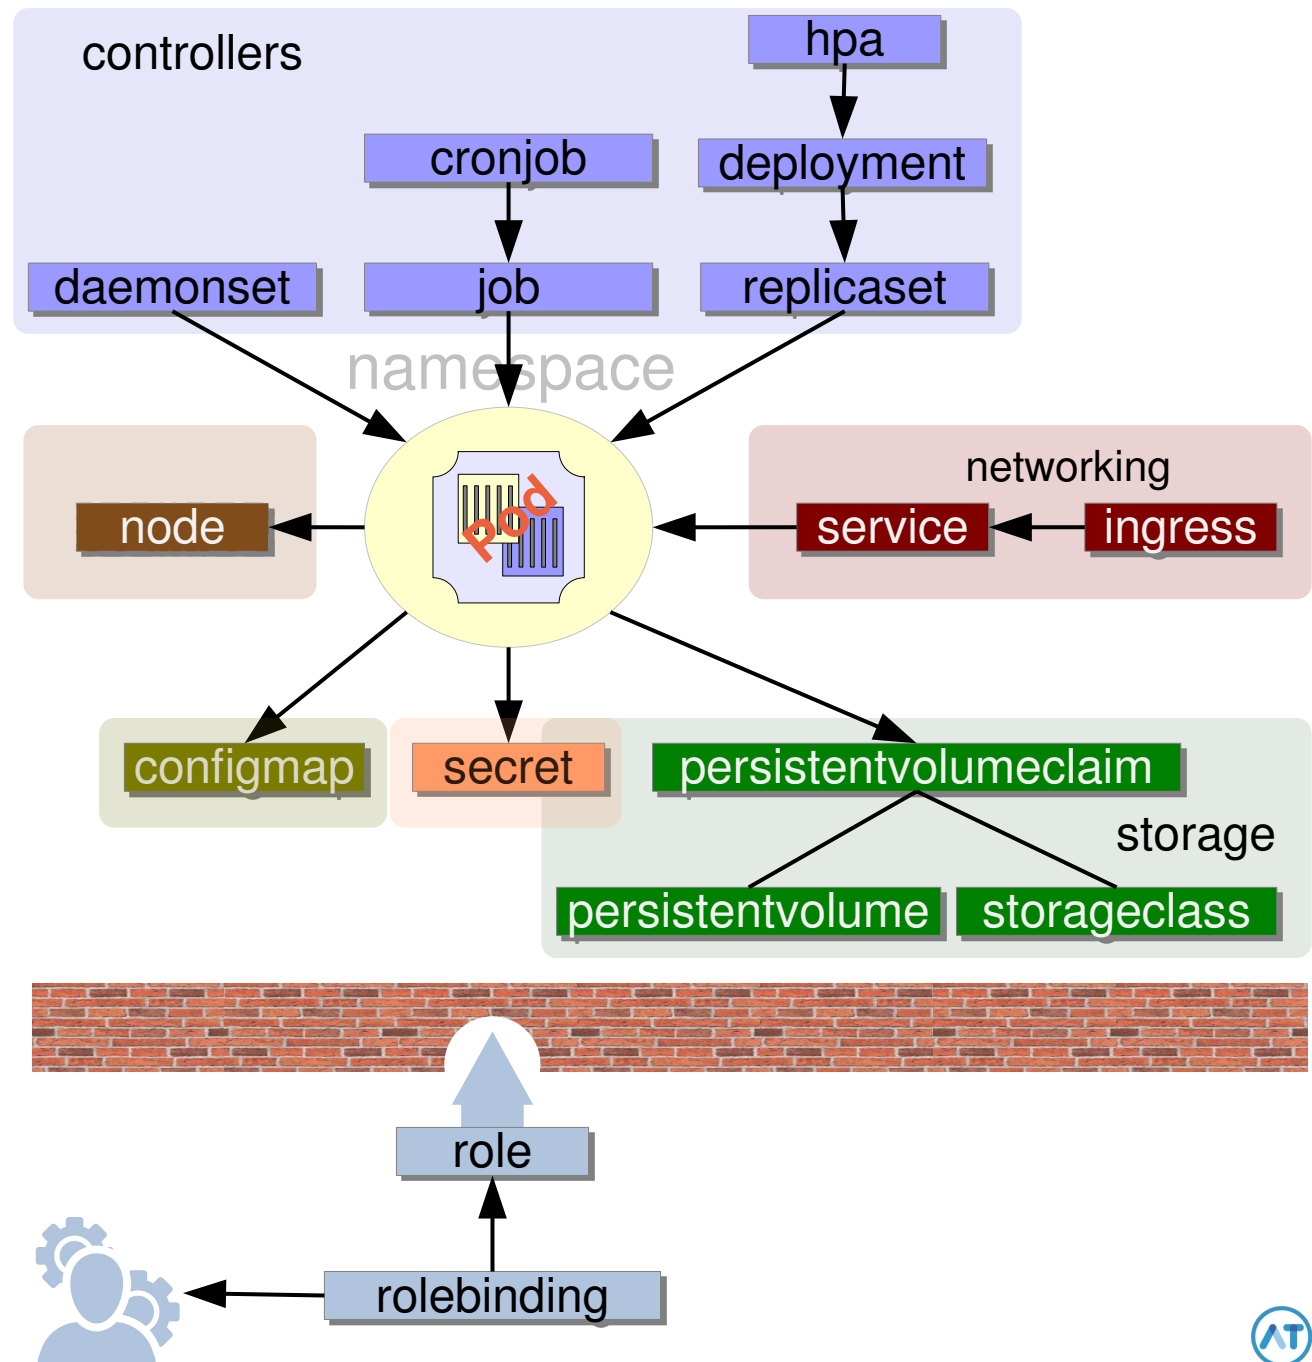
```
apiVersion: v1
kind: Pod
metadata:
  name: apatest
spec:
  containers:
  - name: apacont
    image: atcomp/apachetest
    ports:
    - containerPort: 80
  restartPolicy: Never
```

apa.yml

Kubernetes objects

Kubernetes concept

- numerous object types, like *pod*, *node*, *service*, *deployment*,
- every object has
 - type ('kind')
 - unique name
- object refers to other objects by using
 - labels
(most references *to* pods)
or
 - object names
(most references *from* pods)



Labels

Object labels

- every object has unique name
 - additionally, *labels* can be assigned to be used for
 - selection on command line with `-l` flag
- ```
$ kubectl get pods -l app=webserver
```
- references between objects  
e.g. to assign **Service** object to **Pod** objects

apa.yml

```
apiVersion: v1
kind: Pod
metadata:
 name: apatest
 labels:
 app: webserver
spec:

```

apa-svc.yml

```
apiVersion: v1
kind: Service
metadata:
 name: webservice
spec:
 ports:
 - port: 80
 selector:
 app: webserver
```

# Namespaces

Namespaces: subdivide cluster into various virtual clusters

- per project, per application, per developer team, per department, per ....

```
$ kubectl create ns develop
namespace/develop created
```

```
$ kubectl get ns
NAME STATUS AGE
default Active 278d
develop Active 14s
....
```

```
$ kubectl apply -f apa.yaml -n develop in namespace develop
```

```
$ kubectl get pods -n develop in namespace develop
NAME READY STATUS RESTARTS AGE
apatest 1/1 Running 0 48s
```

- allow
  - separate scope for object names
  - limitation on resource utilization (cpu, memory, storage, number of pods, ...)

# Kubernetes 101 – an introduction

---




Controllers



# Controller types

## Controllers

- pod ('naked pod', 'bare pod') not self-healing
  - pods usually started under supervision of *controller*
  - controllers are part of *controller manager*
- various controller types, like
  - *ReplicaSet* (rs)
    - maintains pod replicas and restarts failing pods
  - *Deployment* (deploy)
    - maintains pod replicas by using ReplicaSet
    - provides rolling updates and rollbacks



# Deployment – create

## Deployment object

- definition implies **ReplicaSet** and **Pod**
- create deployment

```
$ kubectl apply -f apa-deploy.yml
deployment.apps/apadep created
```

```
$ kubectl get deployment
```

| NAME   | READY | UP-TO-DATE | AVAILABLE | ... |
|--------|-------|------------|-----------|-----|
| apadep | 3/3   | 3          | 3         |     |

```
$ kubectl get rs
```

| NAME              | DESIRED | CURRENT | READY | AGE |
|-------------------|---------|---------|-------|-----|
| apadep-7b6fc56c77 | 3       | 3       | 3     | 18s |

```
$ kubectl get pods
```

| NAME                    | READY | STATUS  | RESTARTS | ... |
|-------------------------|-------|---------|----------|-----|
| apadep-7b6fc56c77-6ldcv | 1/1   | Running | 0        |     |
| apadep-7b6fc56c77-bqxhr | 1/1   | Runn    |          |     |
| apadep-7b6fc56c77-hk7qp | 1/1   | Runn    |          |     |

**Deployment**  
image: apachetest:1.14

**ReplicaSet**  
replicas: 3  
image: apachetest:1.14



apa-deploy.yml

```
apiVersion: apps/v1
kind: Deployment
metadata:
 name: apadep
 labels:
 app: webserver
spec:
 replicas: 3
 selector:
 matchLabels:
 app: apapod
 template:
 metadata:
 labels:
 app: apapod
 spec:
 containers:
 - name: apacont
 image: atcomp/apachetest:1.14
 ports:
 - containerPort: 80
```

pod template



# Deployment – rolling update

## Rolling update

- modify deployment manifest file and apply:  
#replicas, environment vars, image version, ...

```
$ kubectl apply -f apa-deploy.yml
deployment.apps/apadep configured
```

first image  
**apachetest:1.14**  
has been modified to  
**apachetest:1.15**

```
$ kubectl get all -o wide
```

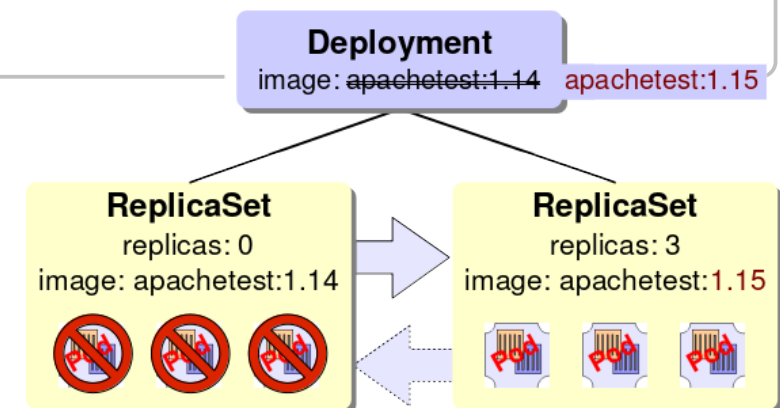
| NAME                   | READY | UP-TO-DATE | AVAILABLE | CONTAINERS | IMAGES                 | ... |
|------------------------|-------|------------|-----------|------------|------------------------|-----|
| deployment.apps/apadep | 3/3   | 3          | 3         | apacont    | atcomp/apachetest:1.15 |     |

| NAME                         | DESIRED | CURRENT | READY | CONTAINERS | IMAGES                 | ... |
|------------------------------|---------|---------|-------|------------|------------------------|-----|
| replicaset/apadep-5b4f756b5c | 3       | 3       | 3     | apacont    | atcomp/apachetest:1.15 |     |
| replicaset/apadep-7b6fc56c77 | 0       | 0       | 0     | apacont    | atcomp/apachetest:1.14 |     |

| NAME                         | READY | STATUS  | IP           | ... |
|------------------------------|-------|---------|--------------|-----|
| pod/apadep-5b4f756b5c-5kvrxx | 1/1   | Running | 10.244.2.213 |     |
| pod/apadep-5b4f756b5c-rflpn  | 1/1   | Running | 10.244.2.212 |     |
| pod/apadep-5b4f756b5c-z5zbj  | 1/1   | Running | 10.244.1.89  |     |

- creates new replicaset within deployment
  - pod replicas replaced one-by-one
  - original replicaset preserved for rollback

```
$ kubectl rollout undo deployment/apadep
```



# Kubernetes 101 – an introduction

---




Networking

# Kubernetes networking

## Communication possibilities

- every pod gets unique IP address on pod network, dynamically assigned
- communication possibilities
  - container-to-container in same pod
    - via `localhost` (loopback interface)
  - pod-to-external
    - via Network Address Translation (NAT)
  - pod-to-pod
    - what is IP address of destination pod?
  - external-to-pod
    - what is IP address of destination pod?



# Networking – IP address of pod

## Example: access via dynamic IP address

- create deployment with 2 pod replicas

```
$ kubectl apply -f apa-deploy.yml
deployment.apps/apadep created

$ kubectl get pods -o wide
NAME READY .. IP NODE
apadep-8654d77c94-928tw 1/1 10.244.1.130 hostb
apadep-8654d77c94-ggv8c 1/1 10.244.2.64 hostc

$ curl 10.244.1.130
<h1> Message from container! </h1>
```

- disadvantages
  - no load balancing
  - when pod terminates, it probably gets another IP address after restart

```
$ kubectl delete pod/apadep-....-928tw
$ kubectl get pods -o wide
NAME READY .. IP NODE
apadep-8654d77c94-7br4b 1/1 10.244.1.131 hostb
apadep-8654d77c94-ggv8c 1/1 10.244.2.64 hostc
```


apa-deploy.yml

```
apiVersion: apps/v1
kind: Deployment
metadata:
 name: apadep
 labels:
 app: webserver
spec:
 replicas: 2
 selector:
 matchLabels:
 app: apapod
 template:
 metadata:
 labels:
 app: apapod
 spec:
 containers:
 - name: apacont
 image: atcomp/apachetest:1.15
 ports:
 - containerPort: 80
```

# Networking – services

## Service object

- separate object
- gets *static* virtual IP address, though dynamically assigned
- attaches pods by using selector referring to pod label
- accessibility determined by *type*
  - ClusterIP**: pod-to-pod communication
  - NodePort**: static port on every node in cluster for external access
- load balancing to attached pods



# Networking – setup ClusterIP service

Example: add *internal* service for webserver

- create service referring to label **app=apapod**
- type **ClusterIP** to provide internal access

```
$ kubectl apply -f apa-deploy.yml
deployment.apps/apadep created
```

```
$ kubectl apply -f apa-svc.yml
service/webserv created
```

```
$ kubectl get service
```

NAME	TYPE	CLUSTER-IP	....	PORT(S)	....
<b>webserv</b>	<b>ClusterIP</b>	10.97.15.1		80/TCP	

```
$ kubectl get all
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/apadep	2/2	2	2	10m

NAME	DESIRED	CURRENT	READY	..
replicaset.apps/apadep-8654d77c94	2	2	2	

NAME	READY	STATUS	RESTARTS	...
pod/apadep-8654d77c94-7br4b	1/1	Running	0	
pod/apadep-8654d77c94-ggv8c	1/1	Running	0	

**apa-deploy.yml**

```
apiVersion: apps/v1
kind: Deployment
metadata:
 name: apadep
spec:
 replicas: 2
 selector:

 template:
 metadata:
 labels:
 app: apapod
 spec:

```

**apa-svc.yml**


```
apiVersion: v1
kind: Service
metadata:
 name: webserv
 labels:
 app: apa
spec:
 type: ClusterIP
 ports:
 - port: 80
 targetPort: 80
 protocol: TCP
 selector:
 app: apapod
```



# Networking – service discovery

## Service discovery by other pods

- via internal DNS
- maintains record for every service:  
`service[.ns.svc.cluster.local]`
- example: activate interactive pod (flags `-it`)



```
$ kubectl run mypod -it --restart=Never --image=atcomp/nwubuntu
root@mypod:/# host webserv
webserv.default.svc.cluster.local has address 10.97.15.1

root@mypod:/# curl webserv
<h1> Message from container! </h1>
```

# Networking – setup NodePort service

Example: add *external* service for webserver

- type **NodePort** to provide external access
- creates port on every node in cluster in range 30000-32767
  - can be specified with keyword **nodePort**

```
$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	....	PORT(S)	....
webserv	NodePort	10.97.15.1		80:32123/TCP	

- access from outside cluster

```
anyhost$ curl hostb:32123 or hosta or hostc
<h1> Message from container! </h1>
```

- access from inside cluster (similar to **ClusterIP**)

```
$ kubectl run mypod -it --image=atcomp/nwubuntu
root@mypod:/# curl webserv
<h1> Message from container! </h1>
```

## apa-deploy.yml

```
apiVersion: apps/v1
kind: Deployment
metadata:
 name: apadep
 labels:
 app: apa
spec:
 replicas: 2
 selector:

 template:
 metadata:
 labels:
 app: apapod
 spec:
 containers:

```

## apa-svcn.yml

```
apiVersion: v1
kind: Service
metadata:
 name: webserv
 labels:
 app: apa
spec:
 type: NodePort
 ports:
 - port: 80
 nodePort: 32123
 protocol: TCP
 selector:
 app: apapod
```

# Kubernetes 101 – an introduction

---




Storage

# Kubernetes Volumes

## Storage

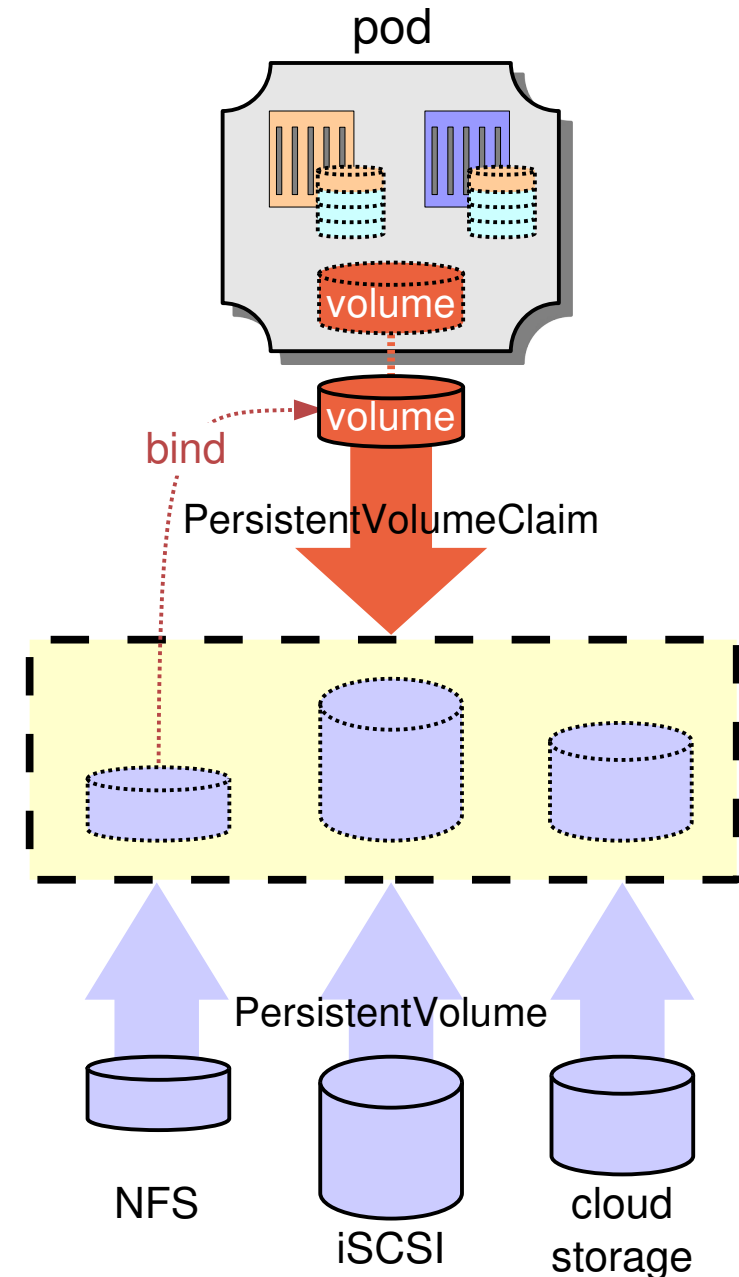
- *stateless pod* – preferred!
  - no need to preserve data
  - non-persistent storage in mini filesystem
- *stateful pod*
  - requires persistent volume
  - various volume types, like
    - local host disk
      - restricted use
    - network storage
      - pod terminated on **hostB** might be restarted on **hostC** and/or
      - pods running on different hosts might share same storage



# Persistent Volumes – static provisioning (1)

## Persistent volumes – *static allocation*

- managed by
  - **PersistentVolume** (PV)
    - *piece of storage* provisioned by administrator
    - example types: NFS, iSCSI
  - **PersistentVolumeClaim** (PVC)
    - *request for storage* to be mounted in pod
    - specific properties can be defined, like size, access mode, performance, ....
- binding of PVC to PV is 1-to-1 mapping



# Persistent Volumes – static provisioning (2)

## Example persistent volume

- create PV of 1GiB based on NFS

```
$ kubectl apply -f pub-pv.yml
persistentvolume/pub-pv created

$ kubectl get pv
NAME CAPACITY ACCESS ... STATUS CLAIM
pub-pv 1Gi RWX Available
```

```
pub-pv.yml
apiVersion: v1
kind: PersistentVolume
metadata:
 name: pub-pv
spec:
 accessModes:
 - ReadWriteMany
 capacity:
 storage: 1Gi
 nfs:
 server: nasi
 path: /nfs/Public
 readOnly: false
```

- request PVC of 500MiB

```
$ kubectl apply -f pub-pvc.yml
persistentvolumeclaim/pub-pvc created

$ kubectl get pvc
NAME STATUS VOLUME CAPACITY ACCESS...
pub-pvc Bound pub-pv 1Gi RWX

$ kubectl get pv
NAME CAPACITY ACCESS ... STATUS CLAIM
pub-pv 1Gi RWX Bound default/pub-pvc
```

```
pub-pvc.yml
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
 name: pub-pvc
spec:
 accessModes:
 - ReadWriteMany
 resources:
 requests:
 storage: 500Mi
```

# Persistent Volumes – static provisioning (3)

## Example persistent volume – cont'd

- create pod using PVC as volume

example:

```
$ kubectl apply -f pub-pod.yml
pod/testpod created
```

```
$ kubectl get pod
```

NAME	READY	STATUS	...
testpod	0/1	Completed	

```
$ kubectl logs testpod
```

```
....
drwxrwxrwx+ ... 4096 Jun 22 2021 Documents
drwxrwxrwx+ ... 4096 Feb 11 2020 Music
drwxrwxrwx+ ... 4096 Dec 17 2020 Photos
```

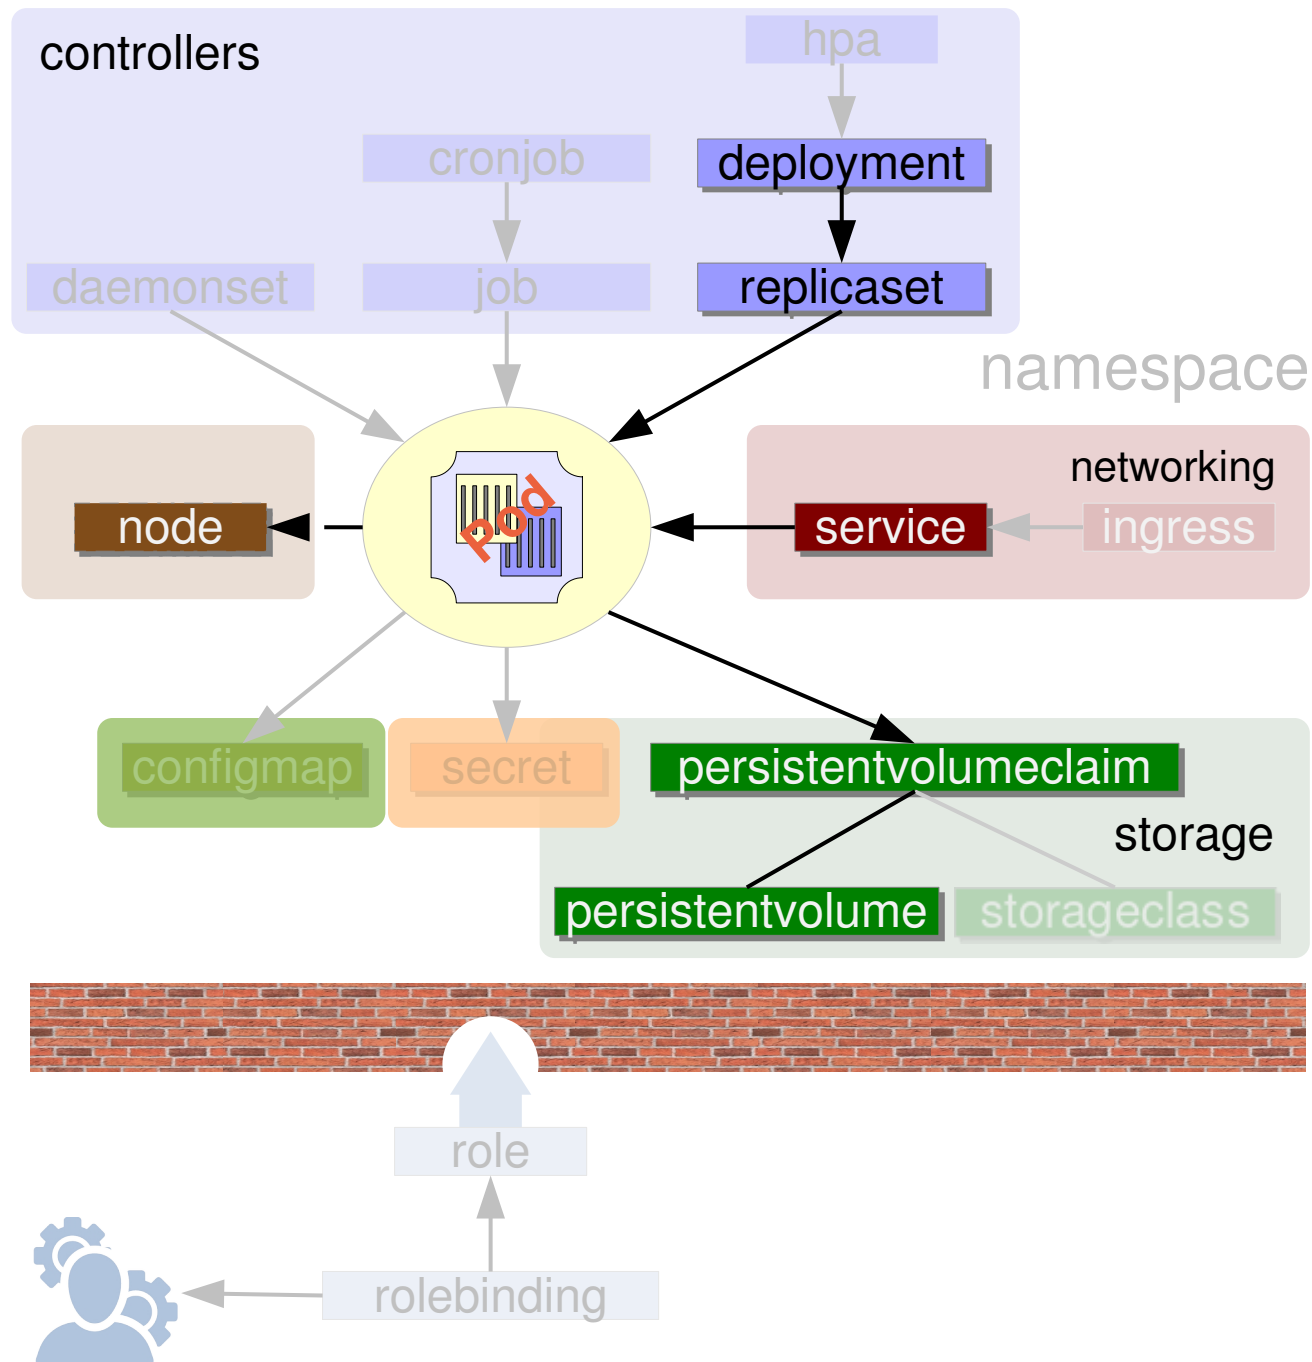
output of  
ls -l /public

```
pub-pod.yml
apiVersion: v1
kind: Pod
metadata:
 name: pub-pod
spec:
 containers:
 - name: lspub
 image: ubuntu
 command: ["ls", "-l", "/public"]
 volumeMounts:
 - name: pubstore
 mountPath: /public
 restartPolicy: Never

 volumes:
 - name: pubstore
 persistentVolumeClaim:
 claimName: pub-pvc
```

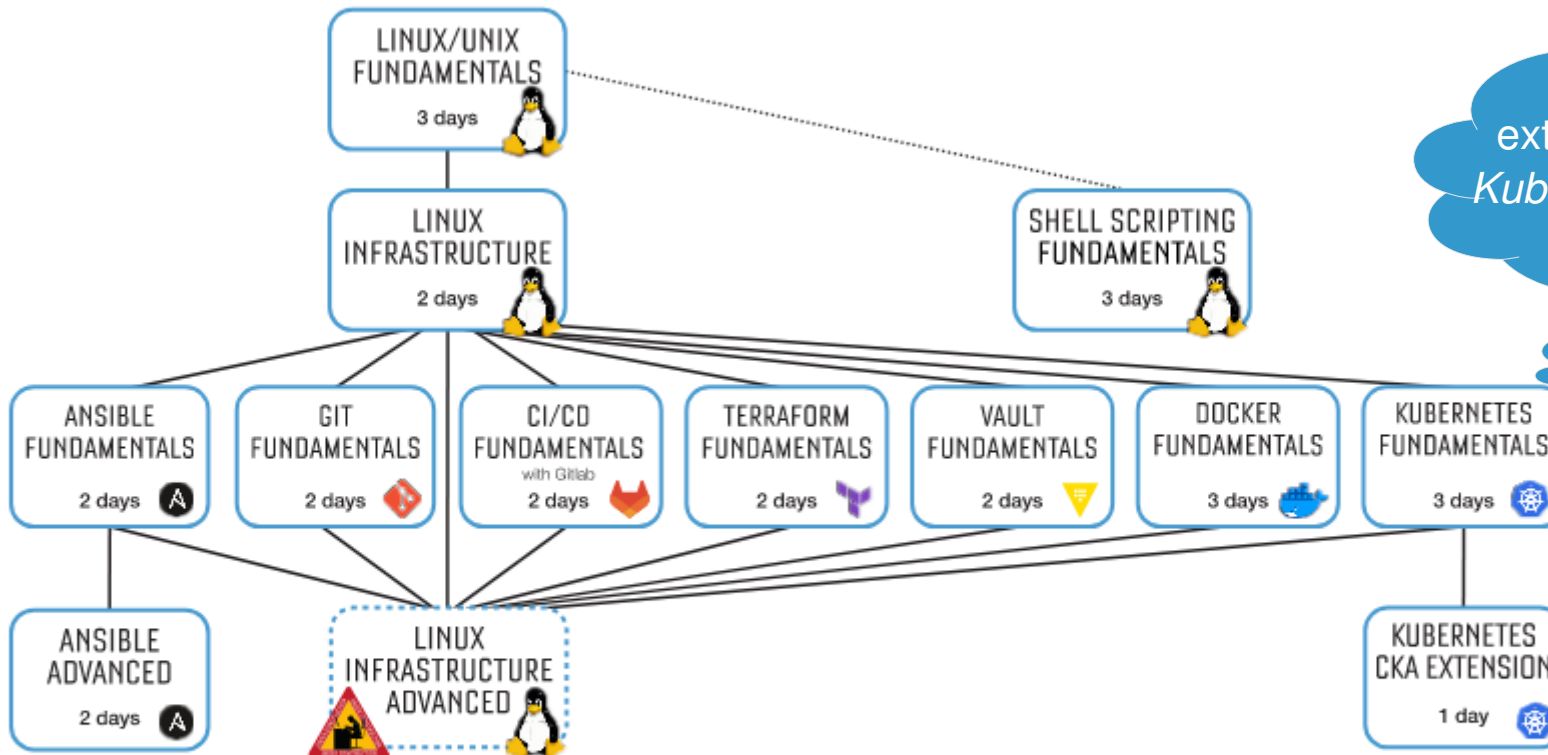
- also possible: *dynamic provisioning*
  - allocate PV dynamically via storage class when creating PVC

# Objects used





# Kubernetes 101 – an introduction



Questions?

```
git clone https://github.com/atcomputing/k8soverview
```